

P8106 Midterm - Code

Kate Colvin (KAC2301), Jeong Yun (Lizy) Choi (JC6452), and Flora Pang (FP2513)

Exploratory Analysis

Loading in Data

```
load("dat1.RData")
load("dat2.RData")

dat1 <- dat1 %>% janitor::clean_names()
dat2 <- dat2 %>% janitor::clean_names()
```

Producing Summary Table

Training and test data have the same distribution of demographic characteristics; there is a difference in time since vaccination and log-transformed antibody levels between training and test data

```
# Combining data for summary table, data cleaning
dat1_com <- dat1 %>% mutate(set = "Training Data")
dat2_com <- dat2 %>% mutate(set = "Testing Data")

dat <- dat1_com %>%
  rbind(dat2_com) %>%
  rename(days_vaccinated = time) %>%
  mutate(race = as.character(race), smoking = as.character(smoking)) %>%
  mutate(race = case_match(
    race, "1" ~ "White", "2" ~ "Asian", "3" ~ "Black", "4" ~ "Hispanic"),
    gender = case_match(gender, 1 ~ "Male", 0 ~ "Female"),
    smoking = case_match(
      smoking, "0" ~ "Never", "1" ~ "Former", "2" ~ "Current"))

# Summary table
dat %>% select(!id) %>%
  tbl_summary(
    by = set,
    label = list(age = "Age", gender = "Gender", race = "Race", smoking = "Smoking",
      height = "Height (cm)", weight = "Weight (kg)", bmi = "BMI",
      diabetes = "Diabetes", hypertension = "Hypertension",
      sbp = "Systolic Blood Pressure (mmHg)", ldl = "LDL Cholesterol (mg/dL)",
      days_vaccinated = "Time Since Vaccinated (days)",
      log_antibody = "Log-Transformed Antibody Level")) %>%
  add_overall() %>% add_p() %>%
  modify_caption("Summary of Patient Testing and Training Data (N=6000)") %>%
  as_gt() %>% tab_options(table.font.size = 10)
```

Table 1: Summary of Patient Testing and Training Data (N=6000)

Characteristic	Overall N = 6,000 ¹	Testing Data N = 1,000 ¹	Training Data N = 5,000 ¹	p-value ²
Age	60.0 (57.0, 63.0)	60.0 (57.0, 63.0)	60.0 (57.0, 63.0)	0.9
Gender				0.7
Female	3,082 (51%)	509 (51%)	2,573 (51%)	
Male	2,918 (49%)	491 (49%)	2,427 (49%)	
Race				0.6
Asian	333 (5.6%)	55 (5.5%)	278 (5.6%)	
Black	1,235 (21%)	199 (20%)	1,036 (21%)	
Hispanic	548 (9.1%)	83 (8.3%)	465 (9.3%)	
White	3,884 (65%)	663 (66%)	3,221 (64%)	
Smoking				0.8
Current	589 (9.8%)	103 (10%)	486 (9.7%)	
Former	1,800 (30%)	296 (30%)	1,504 (30%)	
Never	3,611 (60%)	601 (60%)	3,010 (60%)	
Height (cm)	170.1 (166.1, 174.2)	170.2 (166.1, 174.2)	170.1 (166.1, 174.3)	0.7
Weight (kg)	80 (75, 85)	80 (75, 84)	80 (75, 85)	0.8
BMI	27.60 (25.80, 29.50)	27.60 (25.80, 29.60)	27.60 (25.80, 29.50)	0.9
Diabetes	929 (15%)	157 (16%)	772 (15%)	0.8
Hypertension	2,754 (46%)	456 (46%)	2,298 (46%)	0.8
Systolic Blood Pressure (mmHg)	130 (124, 135)	130 (124, 135)	130 (124, 135)	0.3
LDL Cholesterol (mg/dL)	110 (96, 124)	112 (96, 124)	110 (96, 124)	0.4
Time Since Vaccinated (days)	116 (82, 152)	171 (140, 205)	106 (76, 138)	<0.001
Log-Transformed Antibody Level	10.06 (9.65, 10.45)	9.93 (9.50, 10.32)	10.09 (9.68, 10.48)	<0.001

¹Median (Q1, Q3); n (%)²Wilcoxon rank sum test; Pearson's Chi-squared test

Histograms of Differing Variables by Training and Test Set

```

# Antibody level
plot_sets <- dat %>%
  ggplot(aes(x = log_antibody,
             fill = set,
             color = set)) +
  geom_density(alpha = 0.3, linewidth = 1) +
  labs(x = "Log-Transformed Antibody Level",
       y = "Density",
       title = "Figure 1: Distribution of Log-Transformed Antibody Level, by Data Set") +
  theme_minimal()

# Time since vaccination (days)
plot_days <- dat %>%
  ggplot(aes(x = days_vaccinated,
             fill = set,
             color = set)) +
  geom_density(alpha = 0.3, linewidth = 1) +
  labs(x = "Time Since Vaccinated (days)",
       y = "Density",
       title = "Figure 2: Distribution of Days Since Vaccination, by Data Set") +
  theme_minimal()

plot_sets

```

Figure 1: Distribution of Log-Transformed Antibody Level, by Data Set


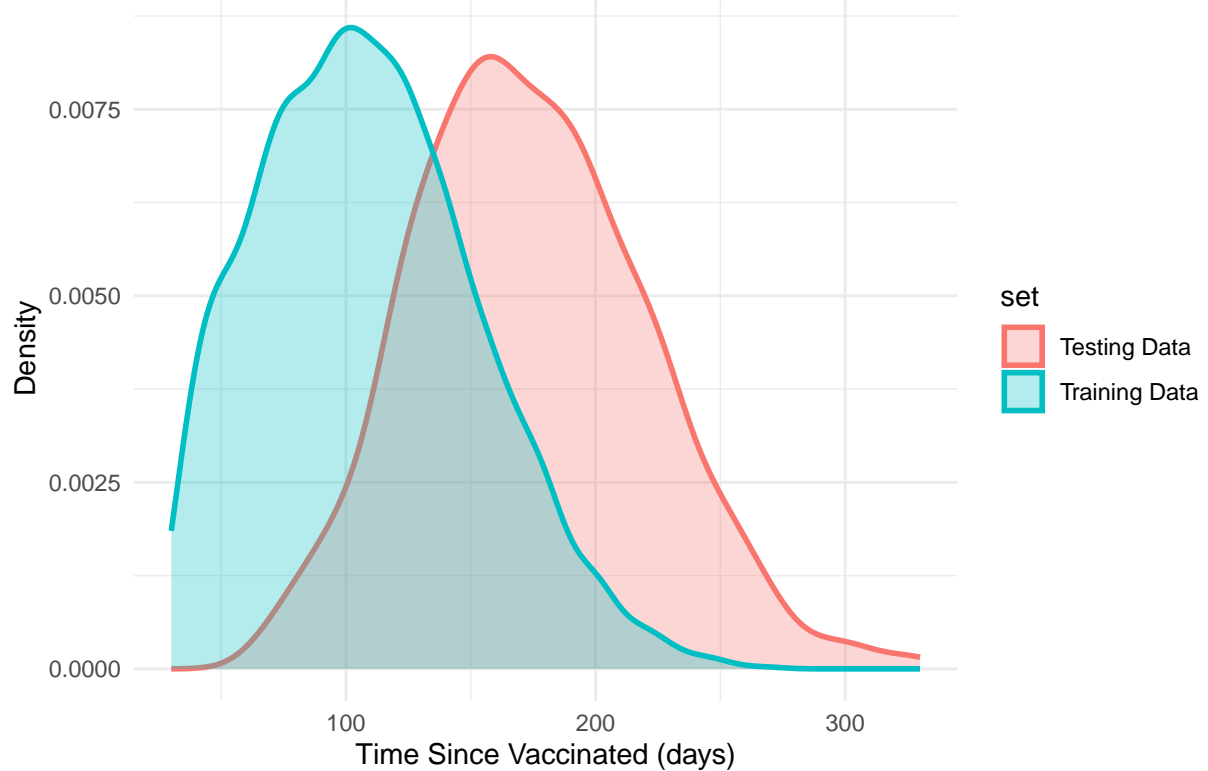



Figure 2: Distribution of Days Since Vaccination, by Data Set



Plots of Log-Transformed Antibody Level, by Categorical Variables

```
# Antibody level, by gender
plot_gender <- dat %>%
  ggplot(aes(x = log_antibody, fill = gender, color = gender)) +
  geom_density(alpha = 0.3, linewidth = 1) +
  labs(x = "Log-Transformed Antibody Level", y = "Density",
       title = "Figure 3: Distribution of Log-Transformed Antibody Level, by Gender") +
  theme_minimal()
plot_gender
```

Figure 3: Distribution of Log-Transformed Antibody Level, by Gender



```
strip_markdown <- function(x) {gsub("\\*\\*", "", x)}

dat %>% select(gender, log_antibody) %>%
  tbl_summary(by = gender) %>% add_p() %>%
  modify_caption("Log-Transformed Antibody Level, by Gender") %>%
  as_kable() %>%
  footnote(general_title = "", general = "Median (Q1, Q3), Wilcoxon Rank Sum Test") %>%
  strip_markdown()
```


Table 2: Log-Transformed Antibody Level, by Gender

Characteristic	Female N = 3,082	Male N = 2,918	p-value
log_antibody	10.20 (9.79, 10.58)	9.93 (9.51, 10.30)	<0.001
Median (Q1, Q3), Wilcoxon Rank Sum Test			

```
# Antibody level, by race
plot_race <- dat %>%
  ggplot(aes(x = log_antibody, fill = race, color = race)) +
  geom_density(alpha = 0.3, linewidth = 1) +
  labs(x = "Log-Transformed Antibody Level",
       y = "Density",
       title = "Figure 4: Distribution of Log-Transformed Antibody Level, by Race") +
  theme_minimal()

plot_race
```

Figure 4: Distribution of Log-Transformed Antibody Level, by Race




```
dat %>%
  select(race, log_antibody) %>%
  tbl_summary(by = race) %>%
  add_p() %>%
  modify_caption("Log-Transformed Antibody Level, by Race") %>%
  as_kable() %>%
  footnote(general_title = "",
          general = "Median (Q1, Q3), Kruskal-Wallis Rank Sum Test") %>%
  strip_markdown()
```

Table 3: Log-Transformed Antibody Level, by Race

Characteristic	Asian N = 333	Black N = 1,235	Hispanic N = 548	White N = 3,884	p-value
log_antibody	10.06 (9.62, 10.44)	10.08 (9.65, 10.44)	10.03 (9.61, 10.42)	10.06 (9.65, 10.46)	0.4
Median (Q1, Q3), Kruskal-Wallis Rank Sum Test					

```
# Antibody level, by smoking status
plot_smoking <- dat %>%
  ggplot(aes(x = log_antibody, fill = smoking, color = smoking)) +
  geom_density(alpha = 0.3, linewidth = 1) +
  labs(x = "Log-Transformed Antibody Level",
       y = "Density",
       title = "Figure 5: Distribution of Log-Transformed Antibody Level, by Smoking") +
  theme_minimal()
plot_smoking
```

Figure 5: Distribution of Log-Transformed Antibody Level, by Smoking



```
dat %>% select(smoking, log_antibody) %>%
  tbl_summary(by = smoking) %>%
  add_p() %>%
  modify_caption("Log-Transformed Antibody Level, by Smoking Status") %>%
  as_kable() %>%
  footnote(general_title = "",
          general = "Median (Q1, Q3), Kruskal-Wallis Rank Sum Test") %>%
  strip_markdown()
```

Table 4: Log-Transformed Antibody Level, by Smoking Status


Characteristic	Current N = 589	Former N = 1,800	Never N = 3,611	p-value
log_antibody	9.91 (9.46, 10.28)	10.10 (9.66, 10.48)	10.07 (9.68, 10.46)	<0.001
Median (Q1, Q3), Kruskal-Wallis Rank Sum Test				

Correlation Matrix of Numerical Variables

```
cor_matrix <- dat %>%
  select(age, height, weight, bmi, sbp, ldl, days_vaccinated, log_antibody) %>%
  rename("Age" = age,
         "Height" = height,
         "Weight" = weight,
         "BMI" = bmi,
         "SBP" = sbp,
         "LDL" = ldl,
         "Days Vaccinated" = days_vaccinated,
         "Log(Antibody)" = log_antibody) %>%
  cor()

cor_plot <- corrplot::corrplot(cor_matrix,
  main = "Figure 6: Correlation Matrix of Numerical Variables",
  mar=c(0,0,1,0), cex.main = 1,
  method = "color",
  addCoef.col = "black",
  tl.col = "black",
  number.cex = 0.8,
  tl.srt = 45,
  order = 'original',
  diag = F)
```

Figure 6: Correlation Matrix of Numerical Variables




Plots of Log-Transformed Antibody Level vs. Selected Numerical Variables

```
# Antibody level vs. BMI
plot_bmi <- dat %>% ggplot(aes(x = bmi, y = log_antibody, fill = set, color = set)) +
  geom_point(alpha = 0.3, size = 2) +
  geom_smooth(method = "lm") +
  labs(y = "Log-Transformed Antibody Level", x = "BMI",
       title = "Figure 7: Log-Transformed Antibody Level vs. BMI") +
  theme_minimal()

# Antibody level vs. Weight
plot_weight <- dat %>%
  ggplot(aes(x = weight, y = log_antibody, fill = set, color = set)) +
  geom_point(alpha = 0.3, size = 2) +
  geom_smooth(method = "lm") +
  labs(y = "Log-Transformed Antibody Level", x = "Weight",
       title = "Figure 8: Log-Transformed Antibody Level vs. Weight") +
  theme_minimal()


# Antibody level vs. Age
plot_age <- dat %>% ggplot(aes(x = age, y = log_antibody, fill = set, color = set)) +
  geom_point(alpha = 0.3, size = 2) + geom_smooth(method = "lm") +
  labs(y = "Log-Transformed Antibody Level", x = "Age",
       title = "Figure 9: Log-Transformed Antibody Level vs. Age") +
  theme_minimal()
```

plot_bmi




plot_weight

Figure 8: Log-Transformed Antibody Level vs. Weight



plot_age

Figure 9: Log-Transformed Antibody Level vs. Age



Model Selection and Training

Since the response variable (log_antibody) is continuous, this project will consider the following models:

- Multiple Linear Regression (MLR) - as a baseline.
- LASSO Regression – to improve predictive performance by selecting important predictors.
- MARS model – allow remain in regression but also capture nonlinear effects

After comparing model performance, the best model will be based on cross-validation results.

Data Pre-processing

```
# Converting categorical variables into factors
dat1 <- dat1 %>%
  mutate(
    gender = factor(gender, levels = c(0, 1), labels = c("Female", "Male")),
    race = factor(race, levels = c(1, 2, 3, 4), labels = c("White", "Asian", "Black", "Hispanic")),
    smoking = factor(smoking, levels = c(0, 1, 2), labels = c("Never", "Former", "Current")),
    diabetes = factor(diabetes),
    hypertension = factor(hypertension)
  )

dat2 <- dat2 %>%
  mutate(
    gender = factor(gender, levels = c(0, 1), labels = c("Female", "Male")),
    race = factor(race, levels = c(1, 2, 3, 4), labels = c("White", "Asian", "Black", "Hispanic")),
    smoking = factor(smoking, levels = c(0, 1, 2), labels = c("Never", "Former", "Current")),
    diabetes = factor(diabetes),
    hypertension = factor(hypertension)
  )

sum(is.na(dat1))

## [1] 0

sum(is.na(dat2))

## [1] 0

dat1 <- dat1 %>% select(-id)
dat2 <- dat2 %>% select(-id)
```

Training multiple linear regression model

```
# train control
set.seed(123)
ctrl1 = trainControl(method = "cv", number = 10)

mlr_model <- lm(log_antibody ~ ., data = dat1)
summary(mlr_model)

##
## Call:
## lm(formula = log_antibody ~ ., data = dat1)
##
## Residuals:
```

```
##      Min      1Q   Median      3Q      Max
## -2.14396 -0.35840  0.02944  0.37802  1.65090
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  26.6751961  2.3149812  11.523 < 2e-16 ***
## age         -0.0205979  0.0019385 -10.626 < 2e-16 ***
## genderMale  -0.2974929  0.0155977 -19.073 < 2e-16 ***
## raceAsian   -0.0060422  0.0344613  -0.175  0.8608
## raceBlack   -0.0075295  0.0196815  -0.383  0.7021
## raceHispanic -0.0417571  0.0273309  -1.528  0.1266
## smokingFormer 0.0219907  0.0173992   1.264  0.2063
## smokingCurrent -0.1934834  0.0269576  -7.177 8.15e-13 ***
## height      -0.0821381  0.0135622  -6.056 1.49e-09 ***
## weight       0.0859034  0.0143481   5.987 2.29e-09 ***
## bmi         -0.2977935  0.0412612  -7.217 6.10e-13 ***
## diabetes1    0.0112795  0.0215643   0.523  0.6010
## hypertension1 -0.0179106  0.0260931  -0.686  0.4925
## sbp          0.0015181  0.0017049   0.890  0.3733
## ldl         -0.0001645  0.0004028  -0.409  0.6829
## time        -0.0003011  0.0001795  -1.677  0.0936 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5503 on 4984 degrees of freedom
## Multiple R-squared:  0.1513, Adjusted R-squared:  0.1488
## F-statistic: 59.25 on 15 and 4984 DF, p-value: < 2.2e-16
```

```
##using caret to see the difference in model result
```

```
mlr_fit <- train(
  log_antibody ~ .,
  data = dat1,
  method = "lm",
  trControl = ctrl1
)

summary(mlr_fit)
```

```
##
## Call:
## lm(formula = .outcome ~ ., data = dat)
##
## Residuals:
##      Min      1Q   Median      3Q      Max
## -2.14396 -0.35840  0.02944  0.37802  1.65090
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  26.6751961  2.3149812  11.523 < 2e-16 ***
## age         -0.0205979  0.0019385 -10.626 < 2e-16 ***
## genderMale  -0.2974929  0.0155977 -19.073 < 2e-16 ***
## raceAsian   -0.0060422  0.0344613  -0.175  0.8608
## raceBlack   -0.0075295  0.0196815  -0.383  0.7021
## raceHispanic -0.0417571  0.0273309  -1.528  0.1266
## smokingFormer 0.0219907  0.0173992   1.264  0.2063
```

```
## smokingCurrent -0.1934834 0.0269576 -7.177 8.15e-13 ***
## height -0.0821381 0.0135622 -6.056 1.49e-09 ***
## weight 0.0859034 0.0143481 5.987 2.29e-09 ***
## bmi -0.2977935 0.0412612 -7.217 6.10e-13 ***
## diabetes1 0.0112795 0.0215643 0.523 0.6010
## hypertension1 -0.0179106 0.0260931 -0.686 0.4925
## sbp 0.0015181 0.0017049 0.890 0.3733
## ldl -0.0001645 0.0004028 -0.409 0.6829
## time -0.0003011 0.0001795 -1.677 0.0936 .
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5503 on 4984 degrees of freedom
## Multiple R-squared: 0.1513, Adjusted R-squared: 0.1488
## F-statistic: 59.25 on 15 and 4984 DF, p-value: < 2.2e-16
```

```
# Evaluating model performance on validation data
mlr_pred <- predict(mlr_model, newdata = dat2)
mlr_rmse <- sqrt(mean((mlr_pred - dat2$log_antibody)^2))
mlr_adj_r2 <- summary(mlr_model)$adj.r.squared
mlr_rmse
```

```
## [1] 0.568318
```

```
mlr_adj_r2
```

```
## [1] 0.1487817
```

```
mlr_caret_pred <- predict(mlr_fit, newdata = dat2)
mlr_caret_rmse <- sqrt(mean((mlr_caret_pred - dat2$log_antibody)^2))
mlr_adj_r2_caret <- summary(mlr_fit)$adj.r.squared
mlr_caret_rmse
```

```
## [1] 0.568318
```

```
mlr_adj_r2_caret
```

```
## [1] 0.1487817
```

To confirm consistency, we also fit the final linear regression model using the caret package in addition to the standard `lm()` function. Both approaches produced identical results in terms of coefficients, RMSE, and adjusted R-squared, confirming that the modeling framework did not influence the outcome. This consistency supports the reliability of our findings regardless of the implementation method.

Training LASSO regression model

Standardizing numerical variables for LASSO using glmnet

```
num_vars <- c("age", "height", "weight", "bmi", "sbp", "ldl", "time") #only continuous variable

preprocess_params <- preProcess(dat1[, num_vars], method = c("center", "scale"))
dat1[, num_vars] <- predict(preprocess_params, dat1[, num_vars])
dat2[, num_vars] <- predict(preprocess_params, dat2[, num_vars])

# Preparing the data matrices for glmnet
x_train <- model.matrix(log_antibody ~ ., dat1)[, -1] # Removing the intercept
y_train <- dat1$log_antibody
```

```

x_valid <- model.matrix(log_antibody ~ ., dat2)[, -1]
y_valid <- dat2$log_antibody

set.seed(123)
lasso_model <- cv.glmnet(x_train, y_train, alpha = 1, standardize = TRUE) # LASSO with cross validation
best_lambda <- lasso_model$lambda.min
lasso_final <- glmnet(x_train, y_train, alpha = 1, lambda = best_lambda, standardize = TRUE) # final model

# predicting with LASSO on validation data
best_lambda

## [1] 1.89747e-05

lasso_pred <- predict(lasso_final, newx = x_valid)
lasso_rmse <- sqrt(mean((lasso_pred - y_valid)^2))
lasso_rmse

## [1] 0.5683693

r2_glmnet <- cor(lasso_pred, y_valid)^2
n <- length(y_valid)
p <- sum(coef(lasso_final, s = "lambda.min") != 0) - 1
lasso_adj_r2 <- 1 - ((1 - r2_glmnet) * (n - 1) / (n - p - 1))

lasso_adj_r2[1]

## [1] 0.1104922

```

using caret package for LASSO to compare the difference between the two methodologies (packages)

```

## using caret for LASSO prediction
set.seed(123)
lasso_fit=train(y=y_train,
               x=x_train,
               method = "glmnet",
               tuneGrid = expand.grid(alpha = 1,
                                     lambda = exp(seq(-10, 1, length = 1000))), # range selected t
               preProcess = c("center", "scale"),
               trControl = ctrl1)

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.

# selected tuning parameter and test error
lasso_caret_best_lambda = lasso_fit$bestTune$lambda
predy2_lasso_caret_fit = predict(lasso_fit, newdata = x_valid)
lasso_caret_rmse = sqrt(mean((y_valid - predy2_lasso_caret_fit)^2))

lasso_caret_best_lambda

## [1] 4.539993e-05

lasso_caret_rmse

## [1] 0.5683938

```

```
test_results <- postResample(pred = predict(lasso_fit, newdata = x_valid),
                             obs = y_valid)
test_results

##      RMSE Rsquared      MAE
## 0.5683938 0.1238137 0.4545446

p_caret <- sum(coef(lasso_fit$finalModel, s = lasso_fit$bestTune$lambda) != 0) - 1
lasso_adj_r2_caret <- 1 - ((1 - test_results["Rsquared"]) * (n - 1) / (n - p_caret - 1))

lasso_adj_r2_caret

## Rsquared
## 0.1104572
```

The LASSO models built using both the glmnet and caret packages produced nearly identical results, with matching RMSE values (0.5684) and very similar adjusted R-squared values (0.0484 vs. 0.0482). This close agreement confirms that the modeling outcome is consistent across both approaches, and that the choice of package did not meaningfully affect the predictive performance.

Training MARS model

```
mars_model <- earth(log_antibody ~ ., data = dat1)
summary(mars_model)

## Call: earth(formula=log_antibody~., data=dat1)
##
##               coefficients
## (Intercept)      10.8474469
## genderMale       -0.2962905
## smokingCurrent   -0.2051269
## h(-0.215005-age)  0.0726888
## h(age- -0.215005) -0.1034029
## h(bmi- -1.46481)  -0.2327459
## h(0.0216075-bmi) -0.1710074
## h(-1.19454-time) -1.4557278
## h(time- -1.19454) -0.0978688
##
## Selected 9 of 10 terms, and 5 of 15 predictors
## Termination condition: RSq changed by less than 0.001 at 10 terms
## Importance: bmi, genderMale, time, age, smokingCurrent, raceAsian-unused, ...
## Number of terms at each degree of interaction: 1 8 (additive model)
## GCV 0.2787787   RSS 1384.431   GRSq 0.2166152   RSq 0.2216218

# predicting with MARS on validation data
mars_pred <- predict(mars_model, newdata = dat2)
mars_rmse <- sqrt(mean((mars_pred - dat2$log_antibody)^2))

mars_rmse

## [1] 0.5327718
```

- The MARS model achieves the lowest RMSE. Therefore MARS will be used as the preferred model for predicting log_antibody. Although further fine-tuning and additional feature exploration could further enhance the model's predictive power.

MARS model tuning


```
tune_grid <- expand.grid(degree = 1:3, nprune = seq(5, 50, by = 5))

mars_tune <- train(log_antibody ~ ., data = dat1, method = "earth",
                  trControl = ctrl1, tuneGrid = tune_grid)

print(mars_tune$bestTune)
```

```
##   nprune degree
##    2      10     1
```

```
ggplot(mars_tune)
```



```
# Train the MARS model with best tuning parameters
mars_model_tune <- train(log_antibody ~ .,
                        data = dat1,
                        method = "earth",
                        trControl = ctrl1,
                        tuneGrid = data.frame(nprune = 10, degree = 1))

summary(mars_model_tune)
```

```
## Call: earth(x=matrix[5000,15], y=c(10.65,9.889,1...), keepxy=TRUE, degree=1,
##           nprune=10)
##
##               coefficients
## (Intercept)    10.8474469
```

```

## genderMale          -0.2962905
## smokingCurrent      -0.2051269
## h(-0.215005-age)    0.0726888
## h(age- -0.215005)   -0.1034029
## h(bmi- -1.46481)    -0.2327459
## h(0.0216075-bmi)    -0.1710074
## h(-1.19454-time)    -1.4557278
## h(time- -1.19454)   -0.0978688
##
## Selected 9 of 10 terms, and 5 of 15 predictors (nprune=10)
## Termination condition: RSq changed by less than 0.001 at 10 terms
## Importance: bmi, genderMale, time, age, smokingCurrent, raceAsian-unused, ...
## Number of terms at each degree of interaction: 1 8 (additive model)
## GCV 0.2787787    RSS 1384.431    GRSq 0.2166152    RSq 0.2216218

print(mars_model_tune$bestTune)

##      nprune degree
## 1      10      1

# Report the final model (regression function)
print(coef(mars_model_tune$finalModel))


##      (Intercept) h(0.0216075-bmi) h(time- -1.19454) h(-1.19454-time)
##      10.84744693      -0.17100738      -0.09786882      -1.45572777
##      genderMale h(age- -0.215005) h(-0.215005-age)      smokingCurrent
##      -0.29629045      -0.10340288      0.07268881      -0.20512685
##      h(bmi- -1.46481)
##      -0.23274594

mars_tune_pred <- predict(mars_model_tune, newdata = dat2)
mars_tune_rmse <- sqrt(mean((mars_tune_pred - dat2$log_antibody)^2))
mars_tune_rmse


## [1] 0.5327718

# variable importance plot for MARS
vip(mars_model_tune$finalModel, type = "nsubsets")

```

```
vip(mars_model_tune$finalModel, type = "rss")
```



```
# Compute R-squared
ss_total_mars <- sum((dat2$log_antibody - mean(dat2$log_antibody))^2)
ss_res_mars <- sum((mars_tune_pred - dat2$log_antibody)^2)
r_squared_mars <- 1 - ss_res_mars / ss_total_mars
```

```
# Get n and number of selected terms (p)
n <- nrow(dat2)
p <- length(mars_model_tune$finalModel$selected.terms)
```

```
# Adjusted R-squared for MARS
mars_adj_r2 <- 1 - ((1 - r_squared_mars) * (n - 1) / (n - p - 1))
```

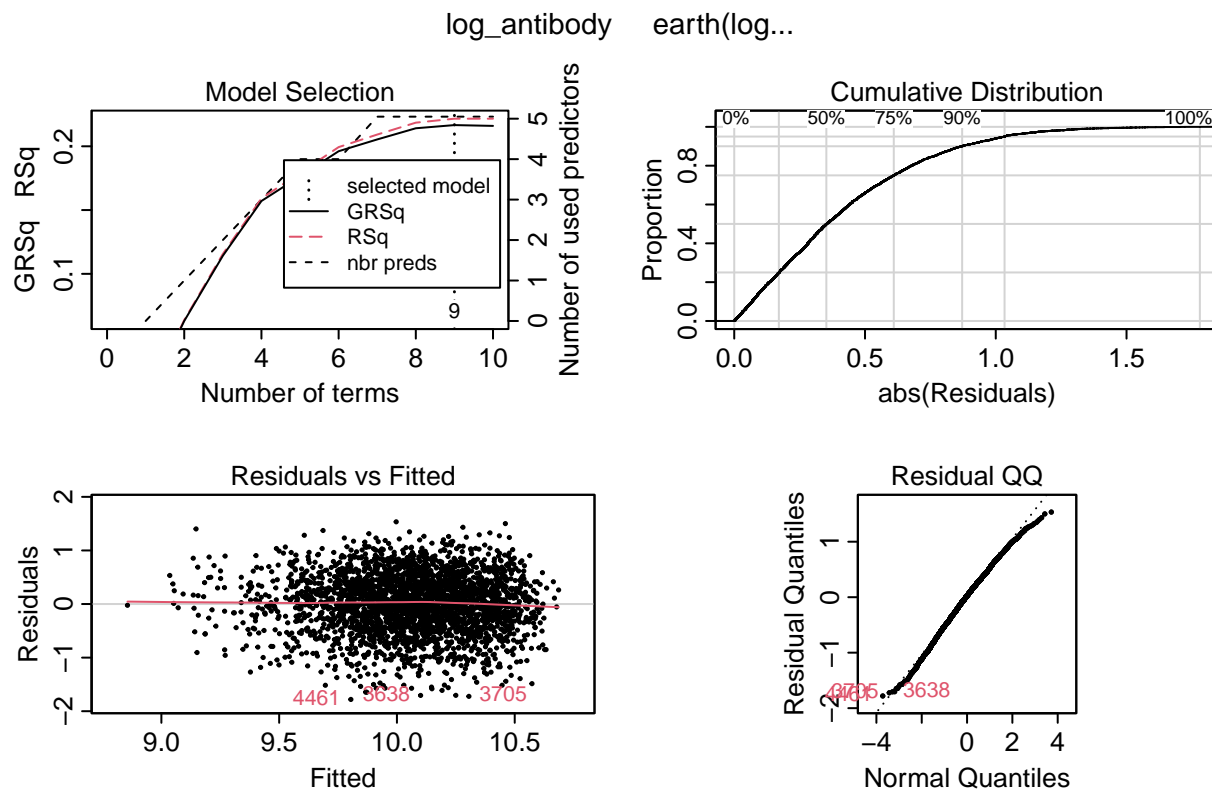
```
mars_adj_r2
```

```
## [1] 0.1689031
```

```
mars_tune_rmse
```

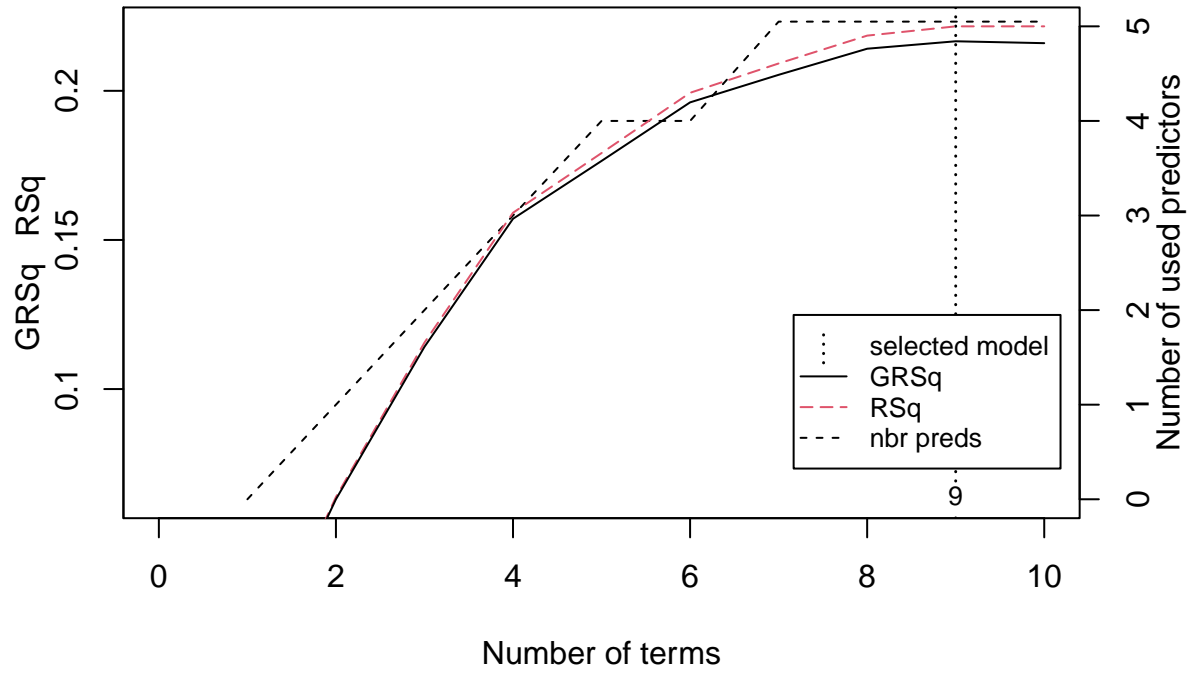
```
## [1] 0.5327718
```

```
plot(mars_model)
```



```
plot(mars_model, which = 1)
```

Model Selection




Results

```
res <- resamples(list(MLR = mlr_fit,
                      LASSO = lasso_fit,
                      MARS = mars_model_tune))
summary(res)

##
## Call:
## summary.resamples(object = res)
##
## Models: MLR, LASSO, MARS
## Number of resamples: 10
##
## MAE
##           Min.    1st Qu.    Median      Mean   3rd Qu.      Max. NA's
## MLR    0.4261808 0.4320149 0.4340128 0.4390390 0.4425761 0.4726409    0
## LASSO  0.4262398 0.4320970 0.4340288 0.4390318 0.4427681 0.4724250    0
## MARS   0.3998101 0.4107959 0.4303208 0.4232841 0.4338504 0.4412629    0
##
## RMSE
##           Min.    1st Qu.    Median      Mean   3rd Qu.      Max. NA's
## MLR    0.5325907 0.5394231 0.5513718 0.5507854 0.5564032 0.5835238    0
## LASSO  0.5328487 0.5395520 0.5514809 0.5508067 0.5562302 0.5833933    0
## MARS   0.5010943 0.5102697 0.5349891 0.5287651 0.5452837 0.5508336    0
##
## Rsquared
##           Min.    1st Qu.    Median      Mean   3rd Qu.      Max. NA's
## MLR    0.1003397 0.1319832 0.1534930 0.1479018 0.1622355 0.1906302    0
## LASSO  0.1000572 0.1325661 0.1528780 0.1478336 0.1616153 0.1911922    0
## MARS   0.1803521 0.1837984 0.2135109 0.2156188 0.2370960 0.2847354    0

# boxplot of RMSE performance for all models
bwplot(res, metric = "RMSE")
```



To ensure consistency across models, we trained the all three model using the caret package as well (model results using different methodologies were comparable) so it would be directly comparable to the others. This allowed us to evaluate all models under the same cross-validation framework and use the `resamples()` function for side-by-side performance comparison. By aligning the training method, we kept the evaluation fair and consistent across approaches.

```
# Summary table
model_perf <- tibble(
  Model = c("Multiple Linear Regression (MLR)",
            "LASSO Regression",
            "MARS (Final, Tuned)"),
  RMSE = c(mlr_rmse, lasso_rmse, mars_tune_rmse),
  `Adjusted R-squared` = c(mlr_adj_r2, lasso_adj_r2, mars_adj_r2),
  Notes = c("Baseline model; assumes linear relationships",
            "Performs variable selection via L1 regularization",
            "Best performance; captures non-linear effects"))

model_perf %>%
  gt() %>%
  tab_header(title = "Table: Model Performance on Independent Test Set (dat2)") %>%
  cols_label(
    Model = "Model",
    RMSE = "RMSE ↓",
    `Adjusted R-squared` = "Adjusted R-squared↑", Notes = "Notes") %>%
  fmt_number(columns = c(RMSE, `Adjusted R-squared`), decimals = 3) %>%
  tab_options(table.font.size = "small")
```

Table: Model Performance on Independent Test Set (dat2)

Model	RMSE ↓	Adjusted R-squared↑	Notes
Multiple Linear Regression (MLR)	0.568	0.149	Baseline model; assumes linear relationships
LASSO Regression	0.568	0.110	Performs variable selection via L1 regularization
MARS (Final, Tuned)	0.533	0.169	Best performance; captures non-linear effects